

```
In [1]: #Experiment No: 10 - Create a Logistic regression model using housing dataset
```

```
In [2]: # Importing necessary libraries
import pandas as pd
import numpy as np
from sklearn import preprocessing
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [3]: # Setting the visual parameters
plt.rc("font", size=14)
sns.set(style="white")
sns.set(style="whitegrid", color_codes=True)

# Load the training and testing datasets using semicolon as the delimiter
train_df = pd.read_csv('train.csv', delimiter=';')
test_df = pd.read_csv('test.csv', delimiter=';')
```

```
In [4]: # Display the first few rows of the training dataset
print("Training Data:\n", train_df.head())
print("\nTest Data:\n", test_df.head())

# Strip any leading/trailing whitespace from column names
train_df.columns = train_df.columns.str.strip()
test_df.columns = test_df.columns.str.strip()
```

# Training Data:

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no

# Test Data:

	age	job	marital	education	default	balance	housing	loan	\
0	30	unemployed	married	primary	no	1787	no	no	
1	33	services	married	secondary	no	4789	yes	yes	
2	35	management	single	tertiary	no	1350	yes	no	
3	30	management	married	tertiary	no	1476	yes	yes	
4	59	blue-collar	married	secondary	no	0	yes	no	

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	cellular	19	oct	79	1	-1	0	unknown	no
1	cellular	11	may	220	1	339	4	failure	no
2	cellular	16	apr	185	1	330	1	failure	no
3	unknown	3	jun	199	4	-1	0	unknown	no
4	unknown	5	may	226	1	-1	0	unknown	no

```
In [5]: # Check for null values
print("\nMissing values in Training Data:\n", train_df.isnull().sum())
print("\nMissing values in Test Data:\n", test_df.isnull().sum())
```

Missing values in Training Data:

```
age      0
job      0
marital  0
education 0
default  0
balance  0
housing  0
loan     0
contact  0
day      0
month    0
duration 0
campaign 0
pdays   0
previous 0
poutcome 0
y        0
dtype: int64
```

Missing values in Test Data:

```
age      0
job      0
marital  0
education 0
default  0
balance  0
housing  0
loan     0
contact  0
day      0
month    0
duration 0
campaign 0
pdays   0
previous 0
poutcome 0
y        0
dtype: int64
```

```
In [6]: # Encoding categorical variables
# Assuming 'y' is the target variable and its values are 'yes'/'no'
train_df['y'] = train_df['y'].map({'yes': 1, 'no': 0})
test_df['y'] = test_df['y'].map({'yes': 1, 'no': 0})

# Splitting the datasets into features and target variable
X_train = train_df.drop('y', axis=1) # Features for training
y_train = train_df['y']              # Target variable for training

X_test = test_df.drop('y', axis=1)   # Features for testing
y_test = test_df['y']                # Target variable for testing

# Convert categorical columns to dummy variables
X_train = pd.get_dummies(X_train, drop_first=True)
X_test = pd.get_dummies(X_test, drop_first=True)
```

```
In [7]: # Align the test set with the train set
X_test = X_test.reindex(columns=X_train.columns, fill_value=0)
```

```
In [8]: # Standardizing the features
scaler = preprocessing.StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [9]: # Creating the Logistic Regression model
model = LogisticRegression(max_iter=1000)
```

```
In [10]: # Fitting the model on the training data
model.fit(X_train, y_train)
```

```
Out[10]: ▼      LogisticRegression ⓘ ⓘ
LogisticRegression(max_iter=1000)
```

```
In [11]: # Making predictions on the test set
y_pred = model.predict(X_test)
```

```
In [12]: # Evaluating the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

```
In [13]: # Displaying the results
print("\nAccuracy:", accuracy)
print("\nConfusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", class_report)
```

Accuracy: 0.9022340190223402

Confusion Matrix:  
[[3905 95]  
 [ 347 174]]

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.98	0.95	4000
1	0.65	0.33	0.44	521
accuracy			0.90	4521
macro avg	0.78	0.66	0.69	4521
weighted avg	0.89	0.90	0.89	4521

```
In [14]: # Visualizing the confusion matrix
plt.figure(figsize=(5, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')  
plt.show()
```

